



# Operating Systems

---

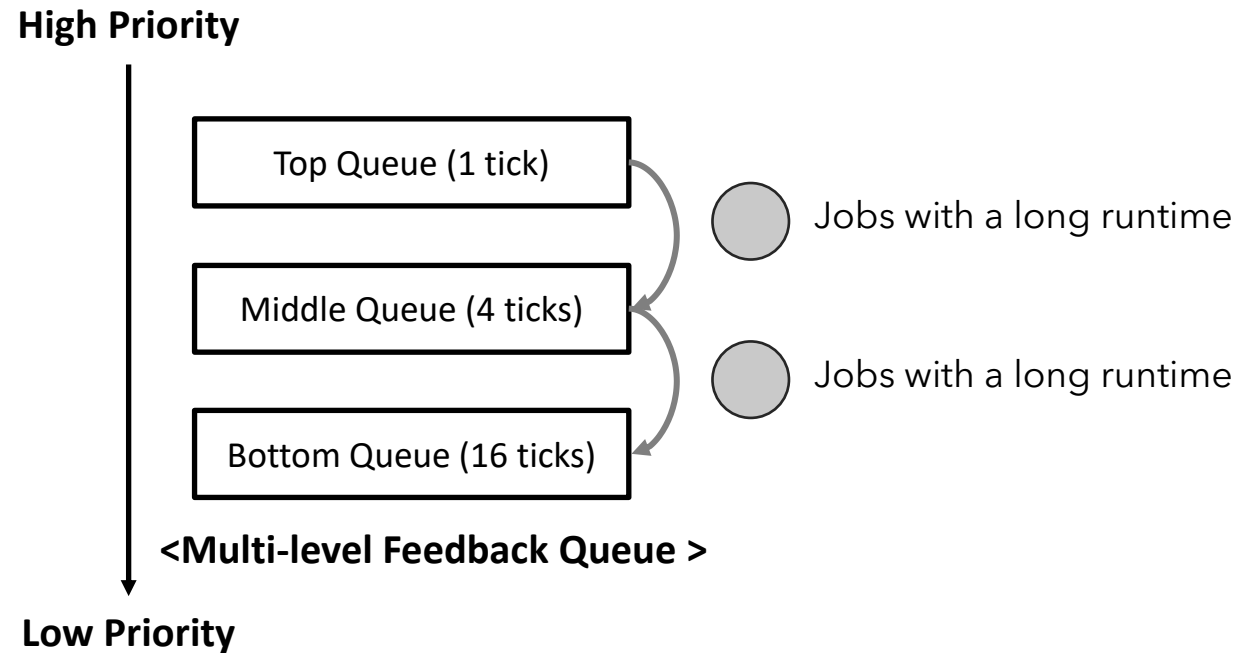
Project I - Scheduler

Minjae Kim

jwya2149@dgist.ac.kr


# Project #1 - Implement xv6 MLFQ scheduler

- Implement a basic MLFQ scheduler in the xv6
  - Modify the default xv6 scheduler to Multi-level Feedback Queue



# Project #1 - Implement xv6 MFLQ scheduler

- What is the default xv6 scheduler?
  - Round-robin scheduler
  - For each timer tick ( $\sim 10\text{ms}$ ), a timer interrupt occurs to incur a context switch
  - Change this round robin to the multi-level feedback!





# Project #1 - Implement xv6 MLFQ scheduler

<default round-robin scheduler code>

- Core function: `void scheduler()` in `proc.c`
- First *for* loop is looping forever
  - This function never returns
  - Find a new process to be scheduled, run it until it yields

**Infinite loop  
(1 tick)**

- Scan `ptable` to find `RUNNABLE` process

```
10 struct {
11     struct spinlock lock;
12     struct proc proc[NPROC];
13 } ptable;
```

- Switch from scheduler to the process
  - After the process yields by timer interrupt, come back to `switch()`

```
322 void scheduler(void) {
323     struct proc *p;
324     struct cpu *c = mycpu();
325     c->proc = 0;
326
327     for(;;){
328         // enable interrupts on this processor.
329         sti();
330
331         // Loop over process table looking for process to run.
332         acquire(&ptable.lock);
333         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
334             if(p->state != RUNNABLE)
335                 continue;
336
337             // Switch to chosen process. It is the process's job
338             // to release ptable.lock and then reacquire it
339             // before jumping back to us.
340             c->proc = p;
341             switchvm(p);
342             p->state = RUNNING;
343
344             switch(&(c->scheduler), p->context);
345             switchkvm();
346
347             // Process is done running for now.
348             // It should have changed its p->state before coming back.
349             c->proc = 0;
350         }
351         release(&ptable.lock);
352     }
353 }
```



# Project #1 - Implement xv6 MLFQ scheduler

- Objectives of this project
  - Understand how context-switches are performed in xv6 code
  - Implement a basic MLFQ scheduler
- Where to look and write code:
  - `proc.c`, `proc.h` (+ etc)

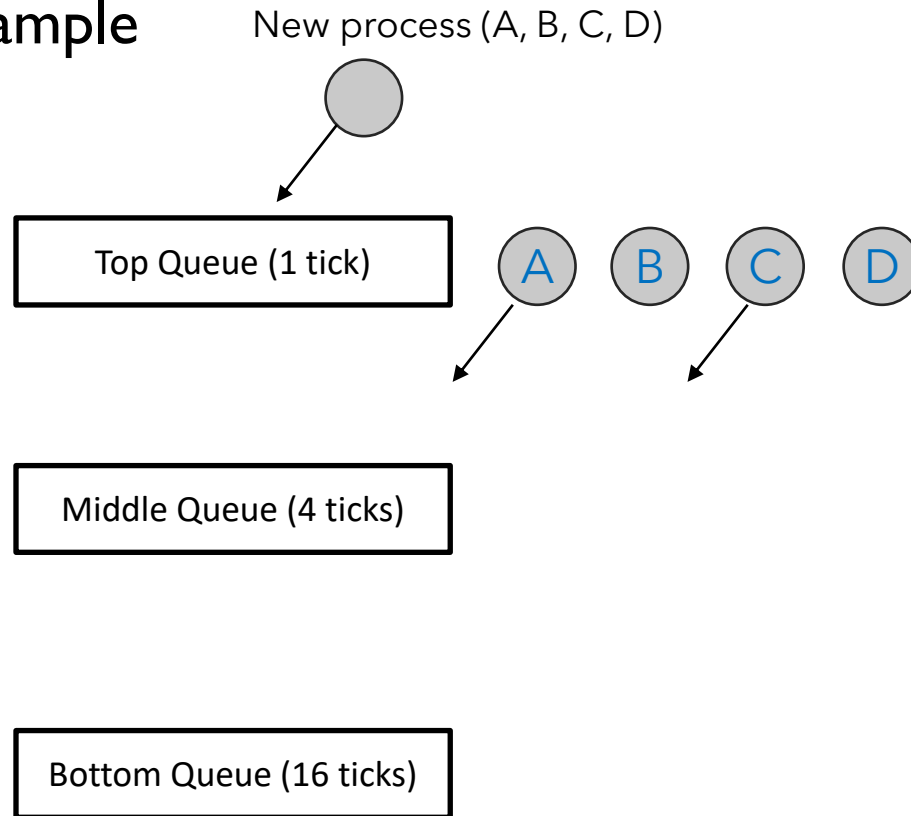
# Project #1 - Implement xv6 MLFQ scheduler

- Rules for MLFQ scheduler

1. There are 3 priority levels – top, middle, bottom
2. Whenever a timer tick interrupt occurs, a process in the top queue is scheduled
  1. Similarly, a process in the top queue is scheduled after previously running process yields the CPU or exits
  2. When a new process is created, it starts from top queue
  3. When scheduling, find a process in the lower queue if there are no processes in the upper queue
3. Processes in the same queue are scheduled by a round-robin policy
  1. However, the time slice of each queue should be considered
  2. Time-slice: # of ticks allowed to each process according to the priority level
4. Time-slice of top/middle/bottom queue is 1 tick / 4 ticks / 16 ticks
5. If a process consumes its time-slice, the process goes to down (except for bottom queue)

# Project #1 - Implement xv6 MLFQ scheduler

- Example




In the top queue, all processes should be scheduled at each timer tick

If a top process does not end in 1 tick, the process goes to the middle queue

# Project #1 - Implement xv6 MLFQ scheduler

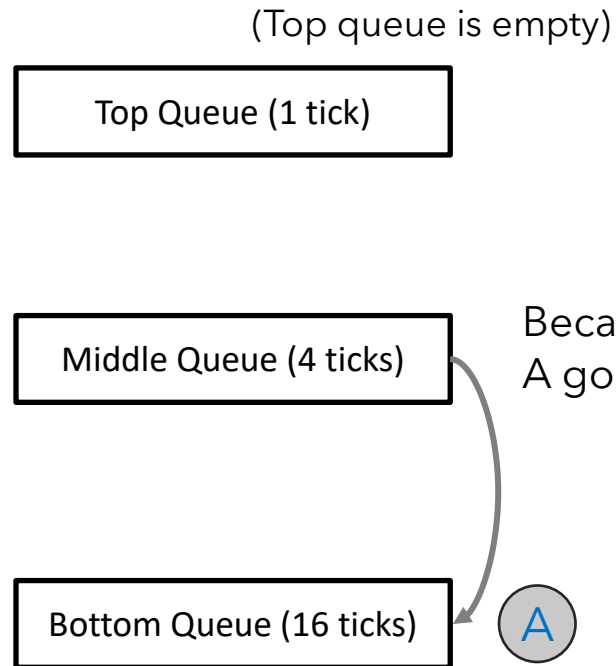
- Example






# Project #1 - Implement xv6 MLFQ scheduler

- Example



Because A doesn't end after using its time-slice,  
A goes to the bottom queue



# Project #1 - Implement xv6 MLFQ scheduler

- Small tips
  - Don't need too many changes in original code – understanding the scheduler flow is important
  - Don't need linked-list queue – recommend to use fixed-sized arrays to represent each priority level
  - To understand more detail about xv6 scheduler, study Chapter 5 in the xv6 book (<https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev11.pdf>)
  - To study MLFQ scheduler, see OSTEP Chapter 8 (<https://pages.cs.wisc.edu/~remzi/OSTEP/cpu-sched.pdf>)
- Build xv6 with 'CPUS = 1' flag to test easier (In Makefile)

```
ifndef CPUS
CPUS := 1
endif
```

# Project #1 - Implement xv6 MLFQ scheduler

- Hints
  - per-process structure: struct proc in proc.h

```
37 // Per-process state
38 struct proc {
39     uint sz;
40     pde_t* pgdir;
41     char *kstack;
42     enum procstate state;
43     int pid;
44     struct proc *parent;
45     struct trapframe *tf;
46     struct context *context;
47     void *chan;
48     int killed;
49     struct file *ofile[NOFILE];
50     struct inode *cwd;
51     char name[16];
52 };
```

- You can add member variables for per-process variable  
e.g.,) allowed ticks, level, ...

# Project #1 - Implement xv6 MFLQ scheduler

- Hints

- New processes are initialized in allocproc()

```
73 static struct proc*
74 allocproc(void)
75 {
76     struct proc *p;
77     char *sp;
78
79     acquire(&ptable.lock);
80
81     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
82         if(p->state == UNUSED)
83             goto found;
84
85     release(&ptable.lock);
86     return 0;
87
88 found:
89     p->state = EMBRYO;
90     p->pid = nextpid++;
91
92     release(&ptable.lock);
93 }
```

- Scan ptable to find an empty space  
(Don't need to be modified)
- Good place to add codes for initialization
  - e.g.,) find empty space in top queue and add p

# Project #1 - Implement xv6 MLFQ scheduler

- Hints

```
322 void scheduler(void) {
323     struct proc *p;
324     struct cpu *c = mycpu();
325     c->proc = 0;
326
327     for(;;){
328         // Enable interrupts on this processor.
329         sti();
330
331         // Loop over process table looking for process to run.
332         acquire(&ptable.lock);
333         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
334             if(p->state != RUNNABLE)
335                 continue;
336
337             // Switch to chosen process. It is the process's job
338             // to release ptable.lock and then reacquire it
339             // before jumping back to us.
340             c->proc = p;
341             switchuvm(p);
342             p->state = RUNNING;
343
344             swtch(&(c->scheduler), p->context);
345             switchkvm();
346
347             // Process is done running for now.
348             // It should have changed its p->state before coming back.
349             c->proc = 0;
350         }
351         release(&ptable.lock);
352     }
353 }
```

- In scheduler(), we need to find which process to schedule
  - Vanilla xv6 scans ptable, but we need to scan queues
  - When scanning queues, we should look at the top level to the bottom level in order
- After swtch(), we can distinguish if the process used just before is over or not
  - We need to modify # of ticks remaining in the process and manage queues according to the case

# Project #1 - Implement xv6 MLFQ scheduler

- Testing

- Run simple bench programs to test whether implemented scheduler works well
  - bench 1~3: fork 10 child processes with a fixed execution time (1/4/16 ticks)
  - bench 4: fork 10 child processes with different execution times ([0, 5, 10, 15, 20, 0, ...] ticks)
  - bench 5: fork 50 child processes with 1 fixed + 49 different times (200 + [0,3,6,9,12,0, ...] ticks)
- To compile bench programs, you need to modify Makefile
  - e.g.,) add `_bench1\`, ... under `UPROGS=`
  - You can make & test any programs if you want
- You need to modify `param.h` to test bench programs
  - `#define NPROC 16` → `64`

```
1 #define NPROC      64 // maximum number of processes
2 #define KSTACKSIZE 4096 // size of per-process kernel stack
3 #define NCPU       8 // maximum number of CPUs
```

```
169 UPROGS=\
170     _cat\
171     _echo\
172     _forktest\
173     _grep\
174     _init\
175     _kill\
176     _ln\
177     _ls\
178     _mkdir\
179     _rm\
180     _sh\
181     _stressfs\
182     _usertests\
183     _wc\
184     _zombie\
185     _mytest\
186     _swtchtest\
187     _realtest\
188     _bench1\
189     _bench2\
```



# Project #1 - Implement xv6 MLFQ scheduler

- Deadline
  - ~ 2022.11.02 (Wed) 23:59
- Hand-in procedure
  - proj1\_201812345.patch
    - Run the following command and upload proj1\_201812345.patch
      - `git diff > proj1_201812345.patch`
    - Check the patch file with Notepad and confirm your modifications are in the patch file
  - Report
    - Submit an 1-2 page report
      - Free format (Korean/English)
      - Description of your implementation
      - Analysis of benchmark programs including comparison with xv6 default scheduler



Finally...

**Do NOT hesitate** to ask questions!

Setup & Mini Projects	Juhyung Park	<a href="mailto:arter97@dgist.ac.kr">arter97@dgist.ac.kr</a>
Project #1	Minjae Kim	<a href="mailto:jwya2149@dgist.ac.kr">jwya2149@dgist.ac.kr</a>
Project #2	Seonggyun Oh	<a href="mailto:sungkyun123@dgist.ac.kr">sungkyun123@dgist.ac.kr</a>
Project #3	Soyoung Han	<a href="mailto:zsally@dgist.ac.kr">zsally@dgist.ac.kr</a>

Thank you!