

Operating Systems Practice

Project #2 – 3-level paging

Seonggyun Oh

(sungkyun123@dgist.ac.kr)



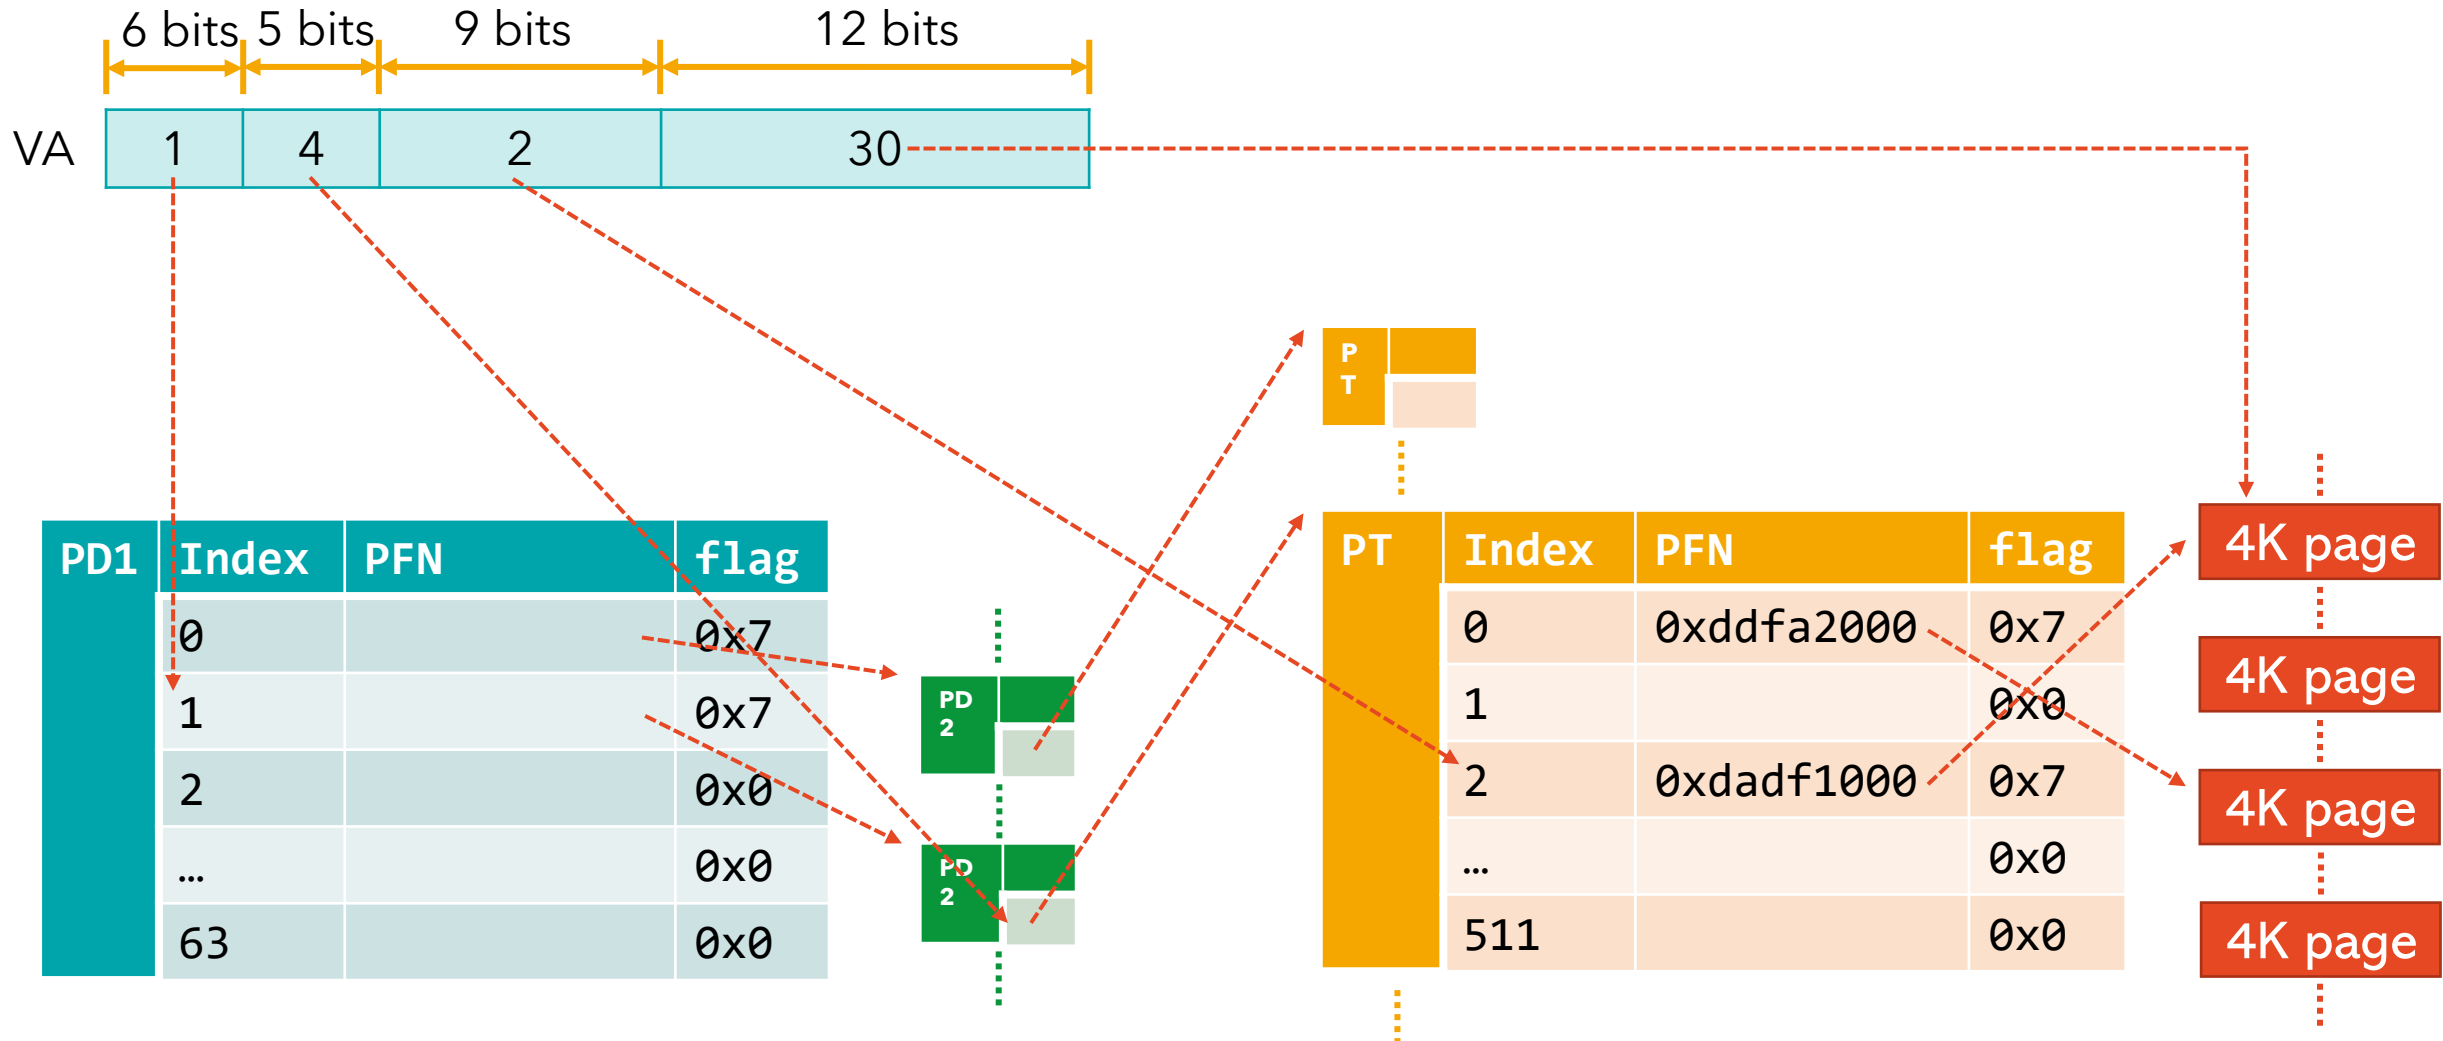


Before we start...

- Clean-up your Git repository for this project
 - `git fetch && git reset --hard && git clean -fdx && git checkout page-counter`
 - Your modifications will be deleted with this command!


3-Level Paging

- 3-level paging



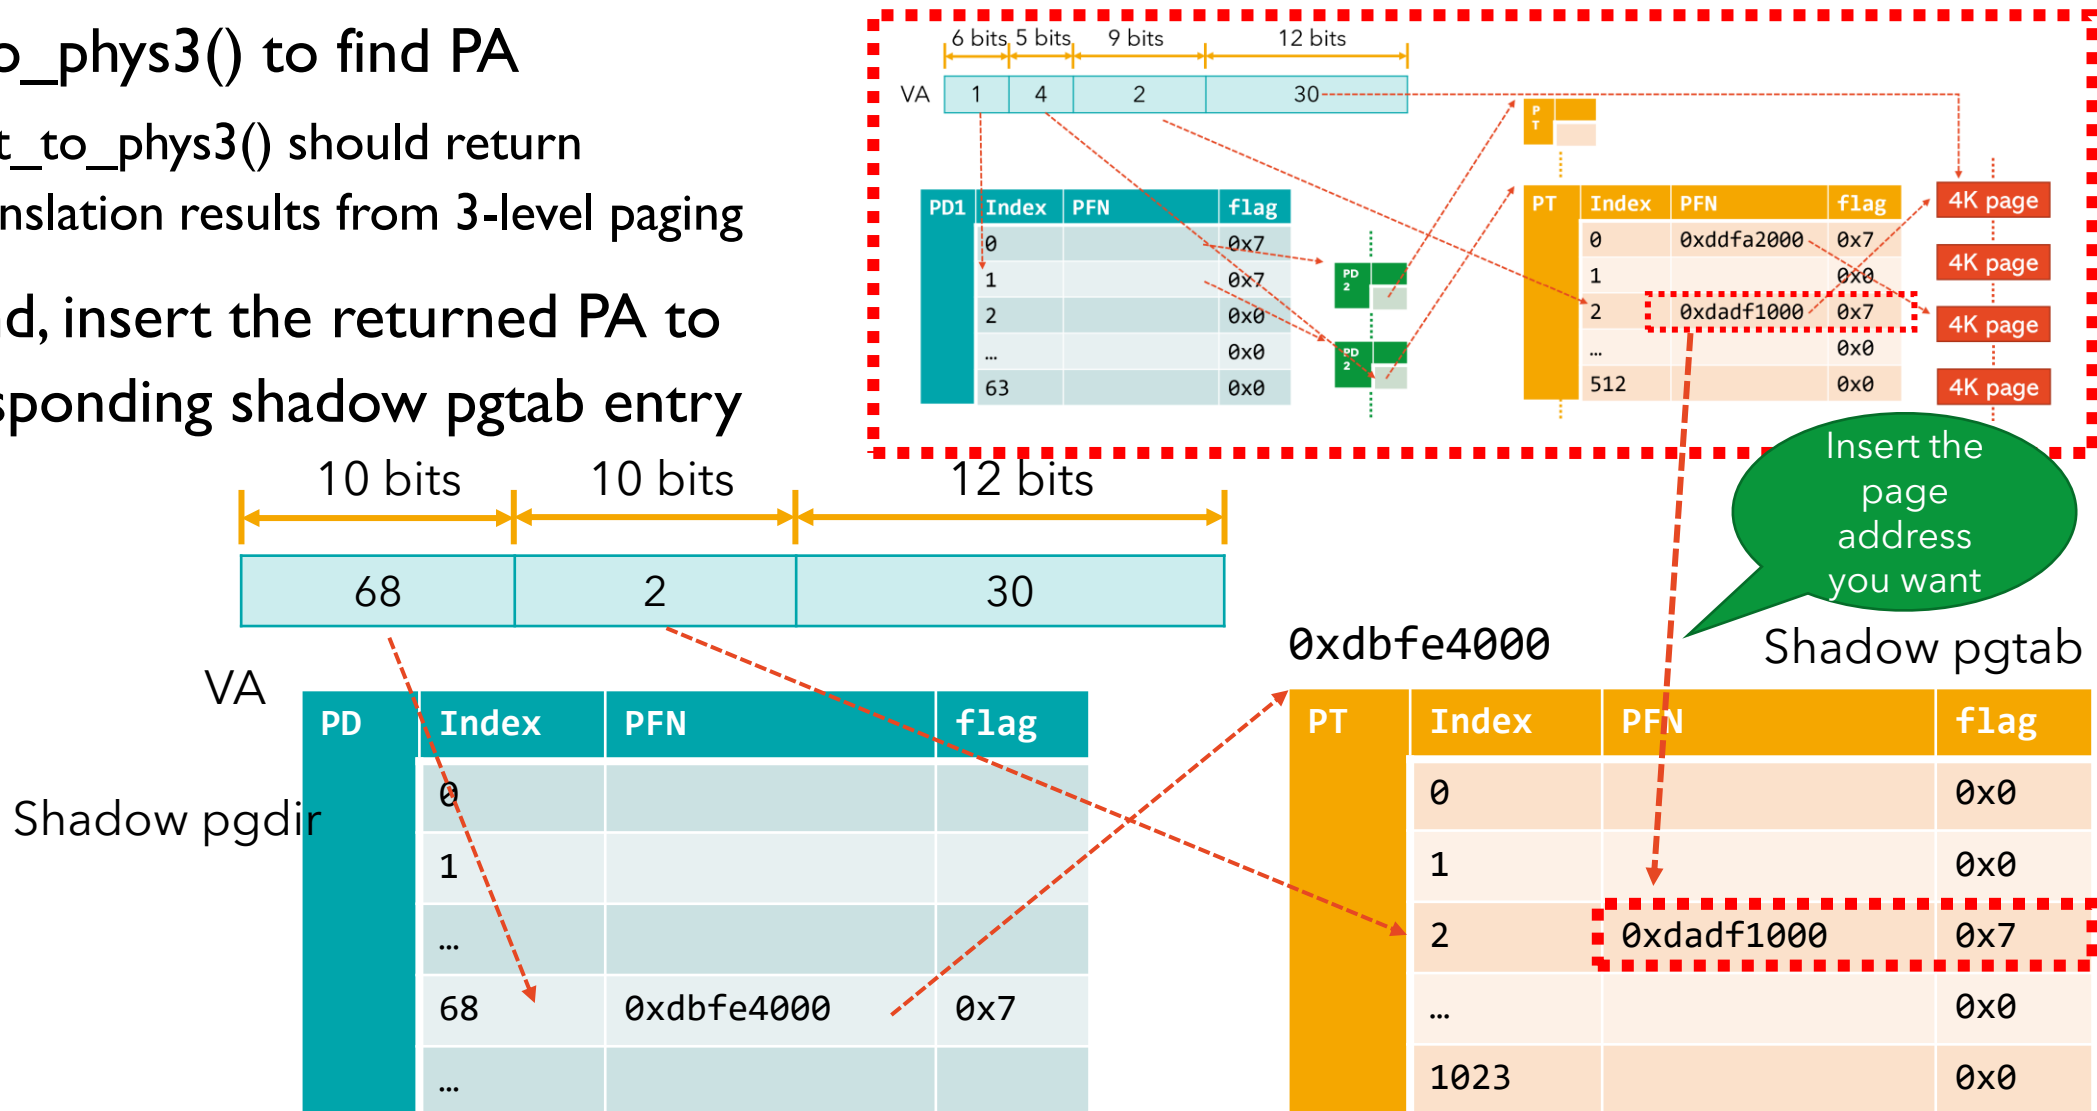
Project #2 – Implementing 3-Level Paging

- *pgdir1*: use upper 6 bits in VA (64 entries)
- *pgdir2*: use 5 bits after 6 bits of *pgdir* in VA (32 entries)
- *pgtab*: use 9 bits after 5 bits of *pgdir2* (512 entries)



Project #2 – Implementing 3-Level Paging

- If page fault occurs, first call `virt_to_phys3()` to find PA
 - `virt_to_phys3()` should return translation results from 3-level paging
- If found, insert the returned PA to corresponding shadow pgtab entry



Project #2 – Implementing 3-Level Paging

- You should modify **allocation, deallocation, traversal code** to support **3-level paging** in context of user's process
 - walkpgdir(), freevm(), deallocvm(), mmu.h
- Use **page-fault concept (software TLB)** to support 3-level paging
 - Modify __virt_to_phys3() and pagefault()
- To evaluate 3-level paging, you should run **usertests** and **memtest**
 - Write the code to output the structure of virtual memory in printvm()
 - Write the code to output the page access count when the process terminates using the leftover space in page directory or page table
- Fill in the code wherever it contains comments written in **//TODO**

Project #2 – Before starting...

- **Because we only consider user processes, you should separate 2-level allocation and 3-level allocation**
 - Kernel's memory paging should be still done in 2-level
 - For 2-level allocation, you should replicate the walkpgdir to `k_walkpgdir()`
 - `shadow_pgdir, kpgdir` uses that function (`k_walkpgdir()`)
 - `walkpgdir()` is only for 3-level allocation and search
 - To do so, we modify `setupkvm()` so that it can distinguish whether the caller is a user process or the kernel
 - `setupkvm(1)` should call `k_walkpgdir()`
 - `setupkvm(0)` should call `walkpgdir()`

Project #2 – Modifying allocating part

- In `mmu.h()`, you should add various **macros** to support **3-level paging**
 - For example, if you extract the index of first page directory entry, make `PDIX` macro that extract the value of upper 6 bits of VA (To do so, the VA must be shifted to 26)
 - Hint
 - Add `PDIX`, `PD2X`, `PTNX`, `PG3ADDR` (or `PG1ADDR` and `PG2ADDR`), `PDIXSHIFT`, `PD2XSHIFT`, `PTNXSHIFT`
- In `walkpgdir()`, you should modify 2-level walking to **3-level walking**
 - Rename original `pgdir` to `pgdir1` and add `pgdir2` to 3-level
 - `pgdir2` should be allocated if it is not allocated yet (like the existing `pgtab`)
 - Return value should be the same as original

Project #2 – Modifying allocating and traversing part

- To track the page access count, you should use margin space in page table
 - We use 9 bits in VA for page table indexing (# of entries per page table: 512 (2^9))
 - xv6 allocate 4 KB pages for page table (# of max entries: 1024 (4KB/4B (sizeof int)))
 - You can use the remaining space (the latter 512 entries) to store page access count
 - When a new page address is faulted (i.e., accessed), increase the count
 - Hint: Write the code in walkpgdir()

Unused!

PT	Index	PFN	flag
	0	0xddfa2000	0x7
	1		0x0
	...		0x0
	511		0x0
	512		
	513		
	...		
	1023		

Project #2 – Modifying deallocating part

- **In freevm(), you should modify 2-level free to 3-level free**
 - freevm() function frees **all allocated page table and page directory**
 - You should free the (new) second page directory too
- **In deallocvm(), you should consider the second page directory**
 - deallocvm() function **frees allocated pages** by walking the page table
 - Originally, the function only considers the first page directory (refer to PGADDR)
 - PGADDR macro changes **the address of input page directory entry** into **the address of next page directory entry**
 - PGADDR only covers 2-level paging, so, you add new macros to support 3-level paging
 - PG3ADDR (or use PG1ADDR and PG2ADDR)
 - If needed, you can make new functions to use in deallocvm()
 - **Hint: You must consider bit-overflow when you traverse valid page to free**
 - **Otherwise, deallocvm() will fall into an endless loop**

Project #2 – Converting VA to PA using Software TLB

- The page fault handler calls `virt_to_phys3()` function to convert VA to PA that is mapped to 3-level paging structure
- You should modify the `__virt_to_phys3()` function to support 3-level paging
- You should modify the `pagefault()` function to support 3-level paging

Project #2 – Evaluating 3-Level Paging

- Make sure that the shadow page table (2-level) returns the same PA as your own 3-level page table
 - `virt_to_phys3(proc->pgdir)` should return the same value as `virt_to_phys2(proc->shadow_pgdir)`!
- Print the page access count when the virtual memory of process is freed
 - Implement the code in `freevm()` function
 - Output

```
va: 0x0, pgtab[0]: 4
va: 0x1000, pgtab[1]: 2
va: 0x2000, pgtab[2]: 5
va: 0x3000, pgtab[3]: 9
va: 0x6000, pgtab[6]: 3
va: 0x7000, pgtab[7]: 3
va: 0xb000, pgtab[11]: 7
va: 0xe000, pgtab[14]: 3
va: 0xf000, pgtab[15]: 3
$
```

Project #2 – Evaluating 3-Level Paging

- To evaluate your implementation of 3-level paging, perform **usertests**
 - If user code may be correct, “ALL TESTS PASSED” will be printed
 - Output

```
Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ usertests
usertests starting
arg test passed
createdelete test
createdelete ok
linkunlink test
```

```
bigdir test
bigdir ok
uio test
pid 551 usertests: trap 13 err 0 on cpu 0 eip 0x3563 addr 0xcf9c--kill proc
uio test done
exec test
ALL TESTS PASSED
$ |
```

Project #2 – Evaluating 3-Level Paging

- Perform **memtest**
 - You must modify **printvm()** function in vm.c to enable to print the information of 3-level
 - You should explain the result in report
 - Output (The output may slightly differ)

```
$ memtest
va: 0x0, pgtab[0]: 3
va: 0x1000, pgtab[1]: 3
va: 0x3000, pgtab[3]: 3
va: 0x4000, pgtab[4]: 3
va: 0xb000, pgtab[11]: 3
initial state of VM of this process

current pid: 3
=Virtual Memory Status=
pgdir: 0x8dd07000
--- 0: pd1e: 0xdc80007, pa: 0xdd07000
----- 0: Page directory2 VA: 0x8dc80000
----- 0: pd2e: 0xdc7f007, pa: 0xdc80000
----- 0: pte: 0xdc81007, pa: 0xdc7f000
----- 2: pte: 0xdc7d007, pa: 0xdc7f000
----- 512: pte: 0x4, pa: 0xdc7f000
----- 514: pte: 0x5, pa: 0xdc7f000

current pid: 3
=Virtual Memory Status=
pgdir: 0x8dd07000
--- 0: pd1e: 0xdc80007, pa: 0xdd07000
----- 0: Page directory2 VA: 0x8dc80000
----- 0: pd2e: 0xdc7f007, pa: 0xdc80000
----- 0: pte: 0xdc81007, pa: 0xdc7f000
----- 2: pte: 0xdc7d007, pa: 0xdc7f000
----- 3: pte: 0xdf9a407, pa: 0xdc7f000
----- 4: pte: 0xde66407, pa: 0xdc7f000
----- 5: pte: 0xde67407, pa: 0xdc7f000
----- 6: pte: 0xdf7b407, pa: 0xdc7f000
----- 7: pte: 0xdf7c407, pa: 0xdc7f000
----- 8: pte: 0xdf7d407, pa: 0xdc7f000
----- 9: pte: 0xdf7e407, pa: 0xdc7f000
----- 10: pte: 0xdf7f407, pa: 0xdc7f000
----- 11: pte: 0xdf80407, pa: 0xdc7f000
----- 12: pte: 0xdf81407, pa: 0xdc7f000
----- 13: pte: 0xdf82407, pa: 0xdc7f000
----- 14: pte: 0xdf83407, pa: 0xdc7f000
----- 15: pte: 0xdf84407, pa: 0xdc7f000
----- 16: pte: 0xdf85407, pa: 0xdc7f000
----- 17: pte: 0xdf86407, pa: 0xdc7f000
----- 18: pte: 0xdf87407, pa: 0xdc7f000
----- 512: pte: 0x4, pa: 0xdc7f000
----- 514: pte: 0x5, pa: 0xdc7f000
----- 523: pte: 0x7, pa: 0xdc7f000
va: 0x0, pgtab[0]: 4
va: 0x1000, pgtab[1]: 2
va: 0x2000, pgtab[2]: 5
va: 0x3000, pgtab[3]: 9
va: 0x6000, pgtab[6]: 3
va: 0x7000, pgtab[7]: 3
va: 0xb000, pgtab[11]: 7
va: 0xe000, pgtab[14]: 3
va: 0xf000, pgtab[15]: 3
$ |
```


Project #2 – 3-Level Paging

- Deadline
 - ~ 2022.11.23 (Wed) 23:59
- Hand-in procedure
 - p2_201812345.patch
 - Run the following command and upload p2_201812345.patch
 - `git diff > p2_201812345.patch`
 - Check the patch file with Notepad and confirm your modifications are in the patch file
 - Report
 - Submit an 1~3 pages report
 - Free format (Korean/English)
 - Description of your implementation in detail (walkpgdir, deallocvm, freevm)
 - Explain your test code and answer the following two tests and explain why the results are as follows
 - Insert test code result image



Extra points

- If you found a mistake in the skeleton code
- If you found a different/better method
- Kernel memory leak test has been disabled in usertests
 - <https://github.com/dgist-datalab/xv6/commit/63776c9cc>
 - If you manage to pass this test (after reverting the commit)
- Write a report for extra points



Finally ...

Do NOT hesitate to ask questions!

Mini Project #1, #2 Juhyung Park

arter97@dgist.ac.kr

Project #1 Minjae Kim

jwya2149@dgist.ac.kr

Project #2 Seonggyun Oh

sungkyun123@dgist.ac.kr

Project #3 Soyoung Han

zsally@dgist.ac.kr



Thank You!