

Operating Systems Practice

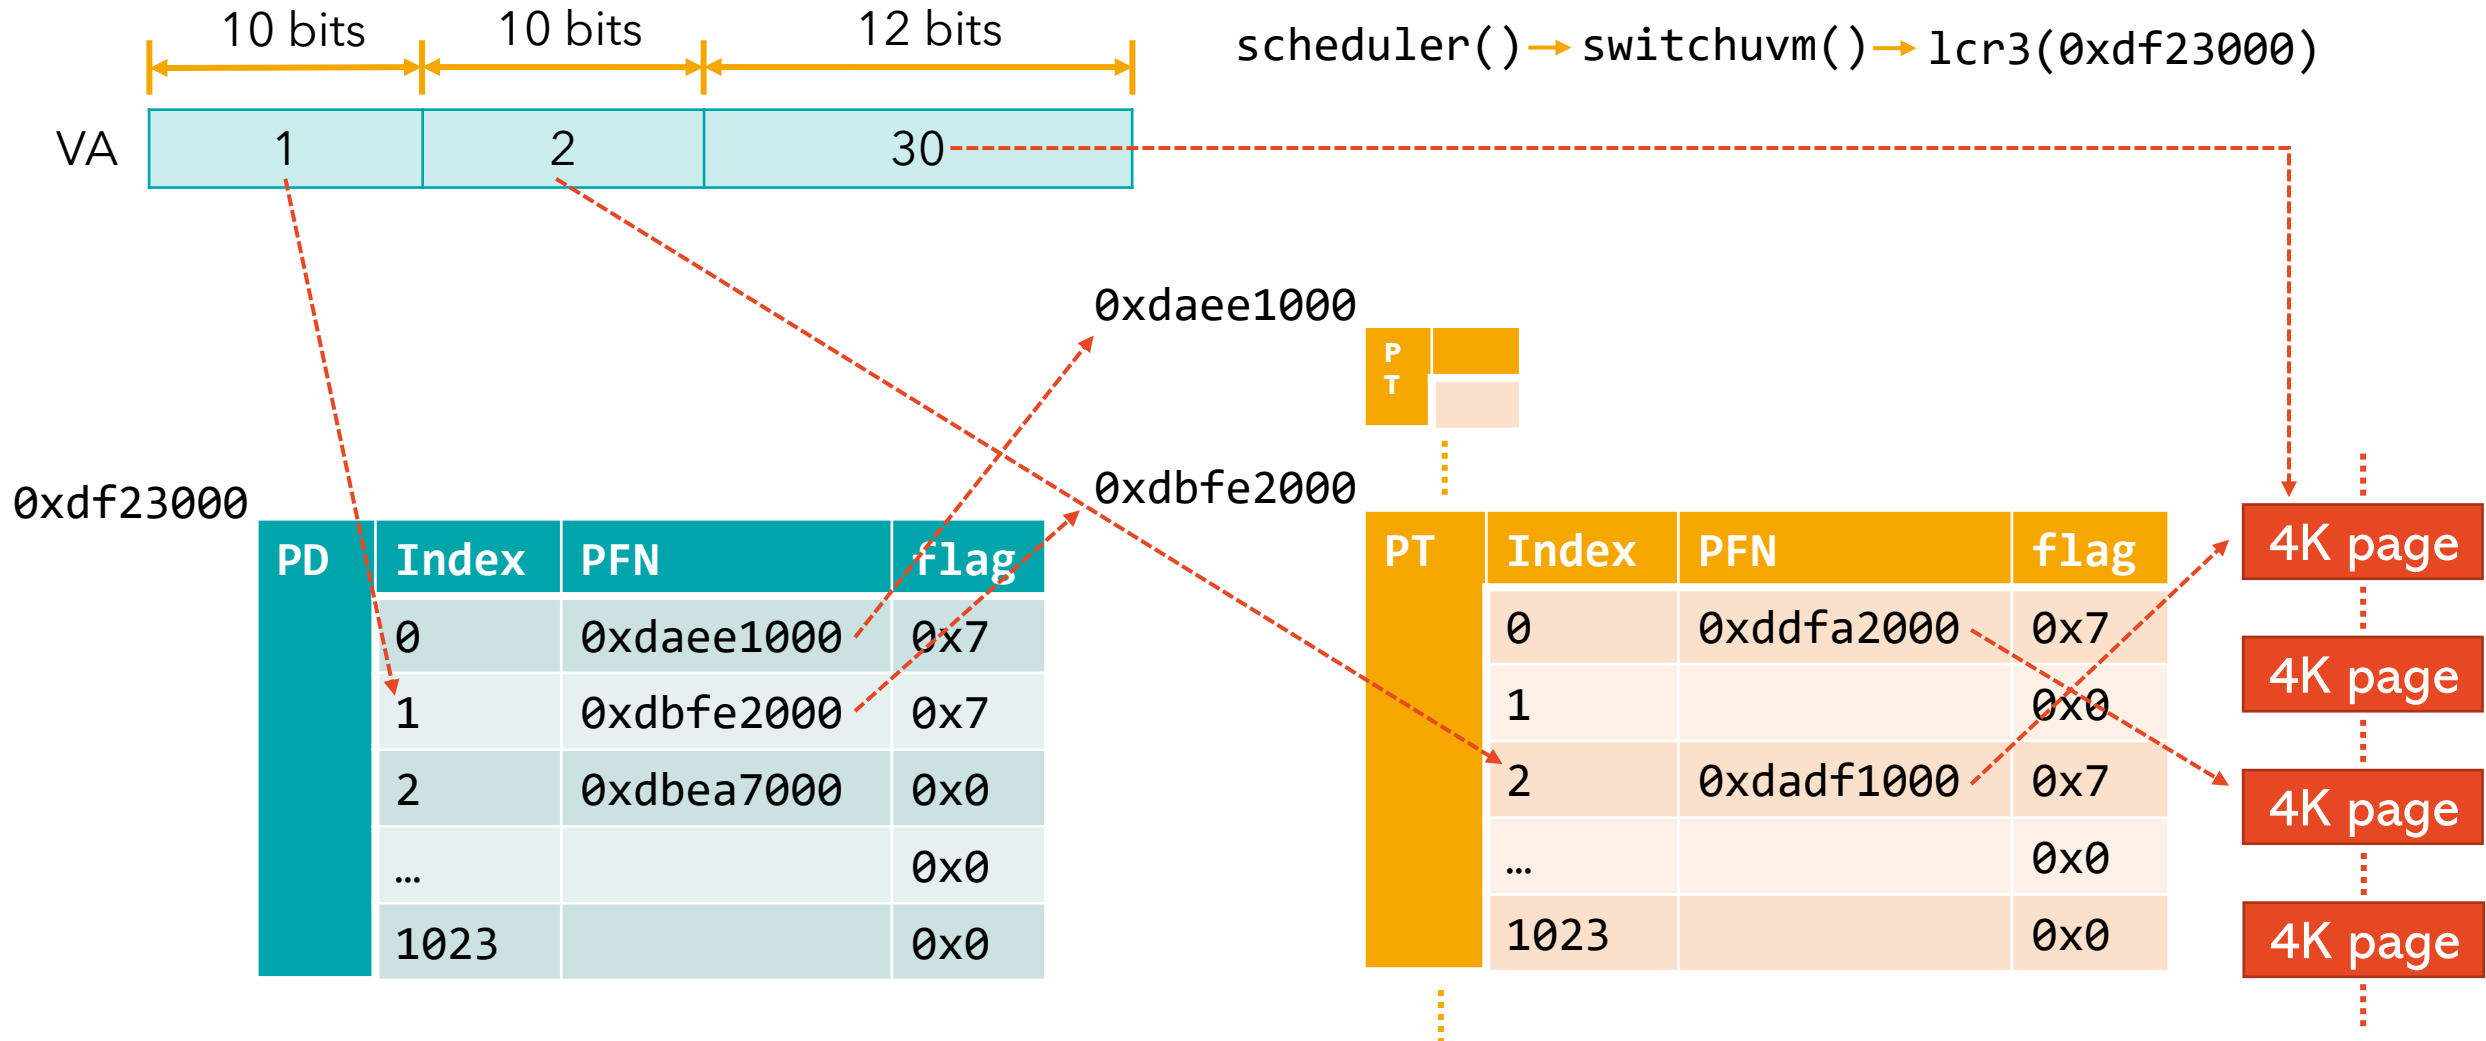
Project #2 - 3-level paging (Part I)

Juhyung Park

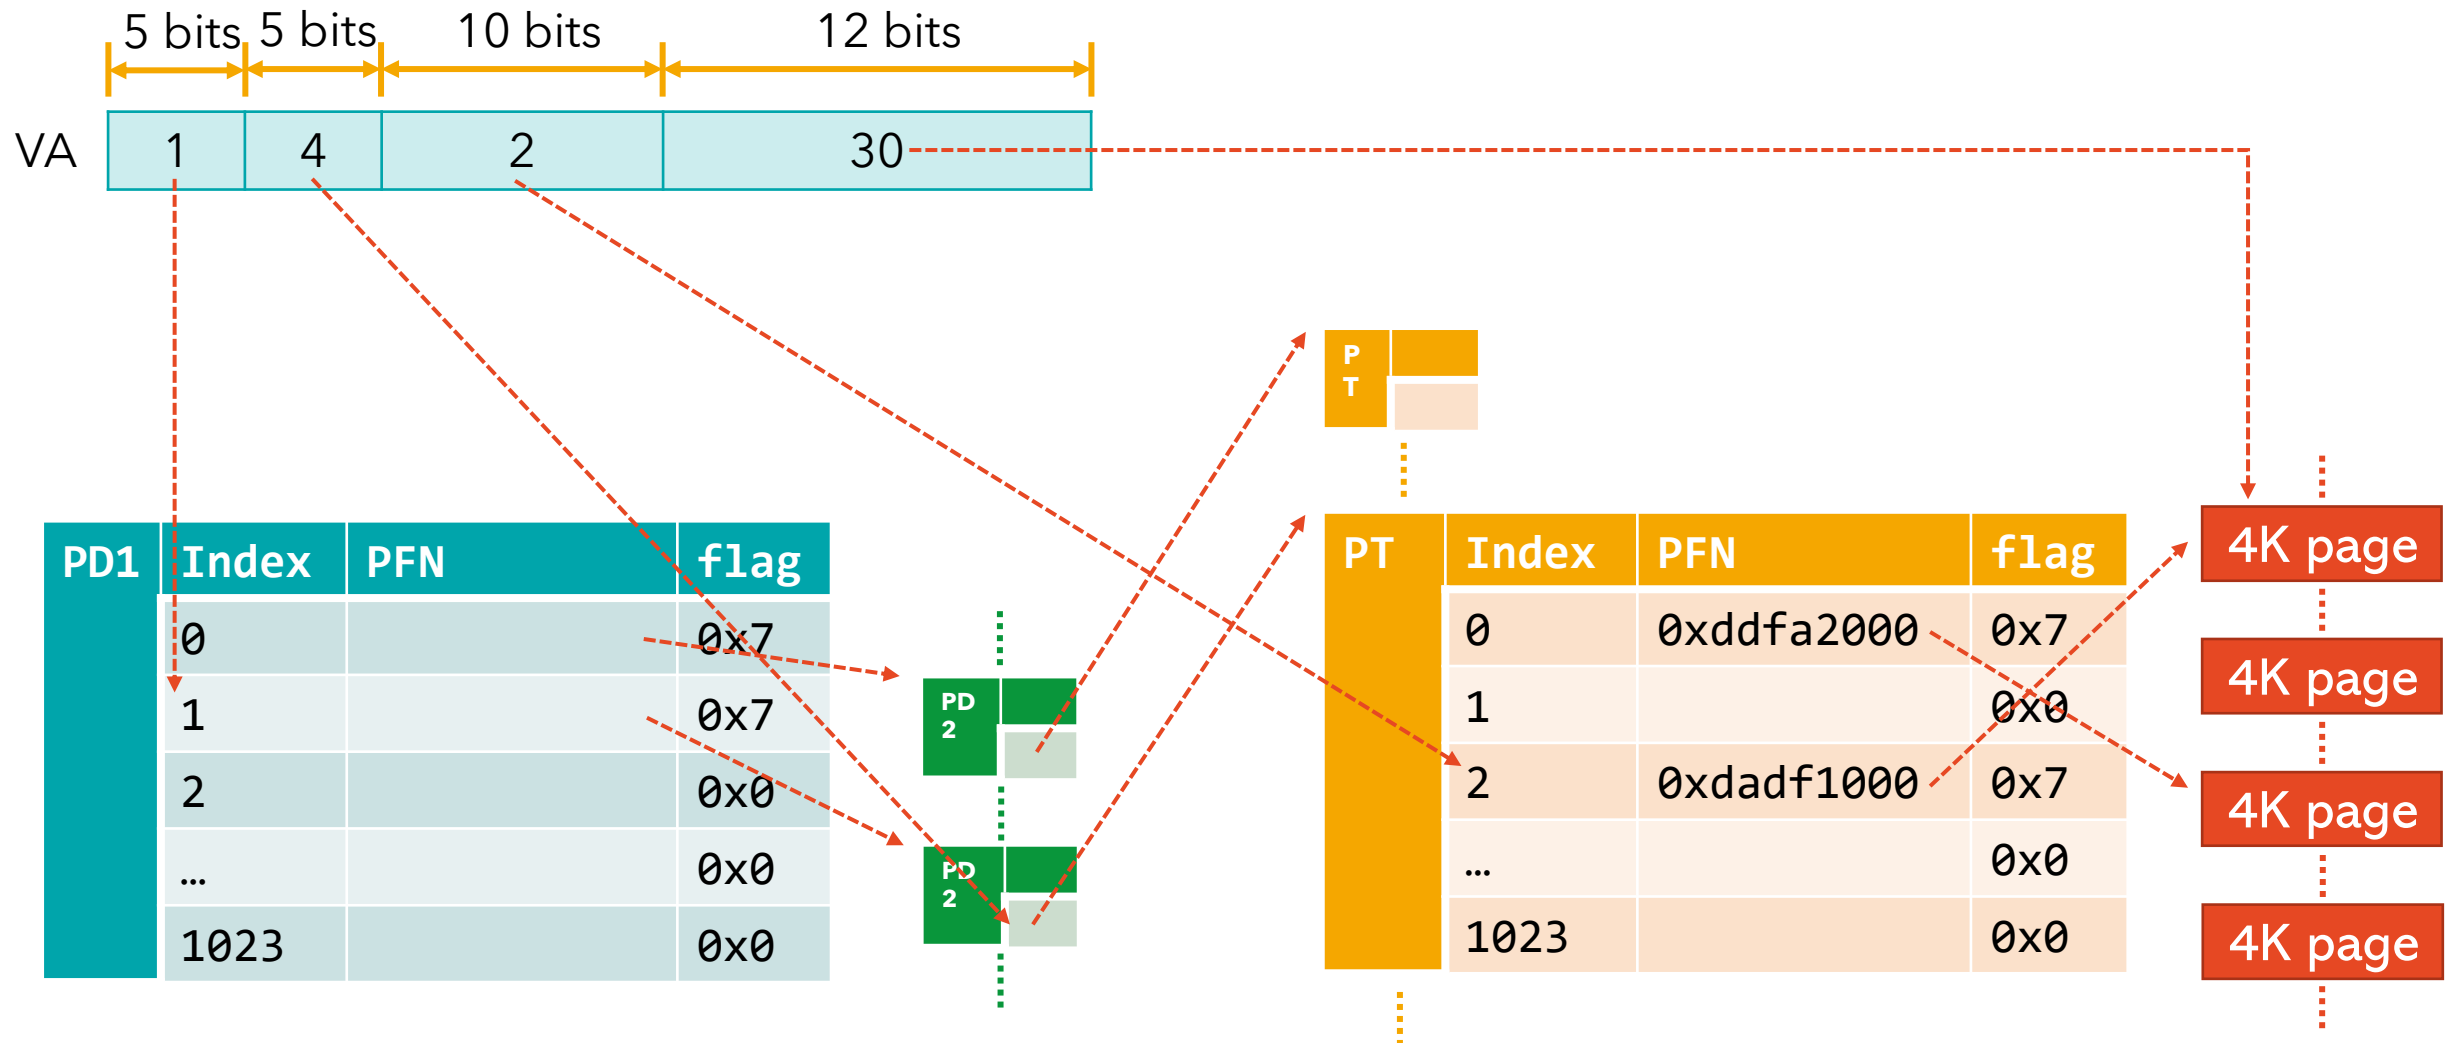
arter97@dgist.ac.kr



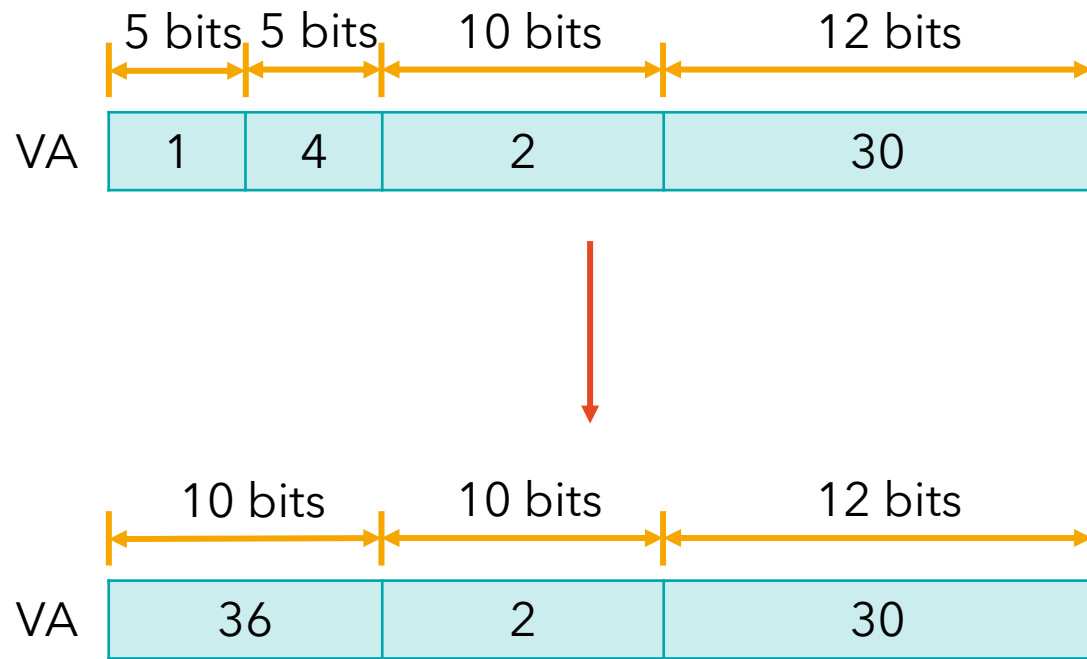
Recall how VA is translated to PA (2-level page table)



How can we customize VA to PA translation? (e.g., 3-level)



x86-32 MMU still expects 2-level VA

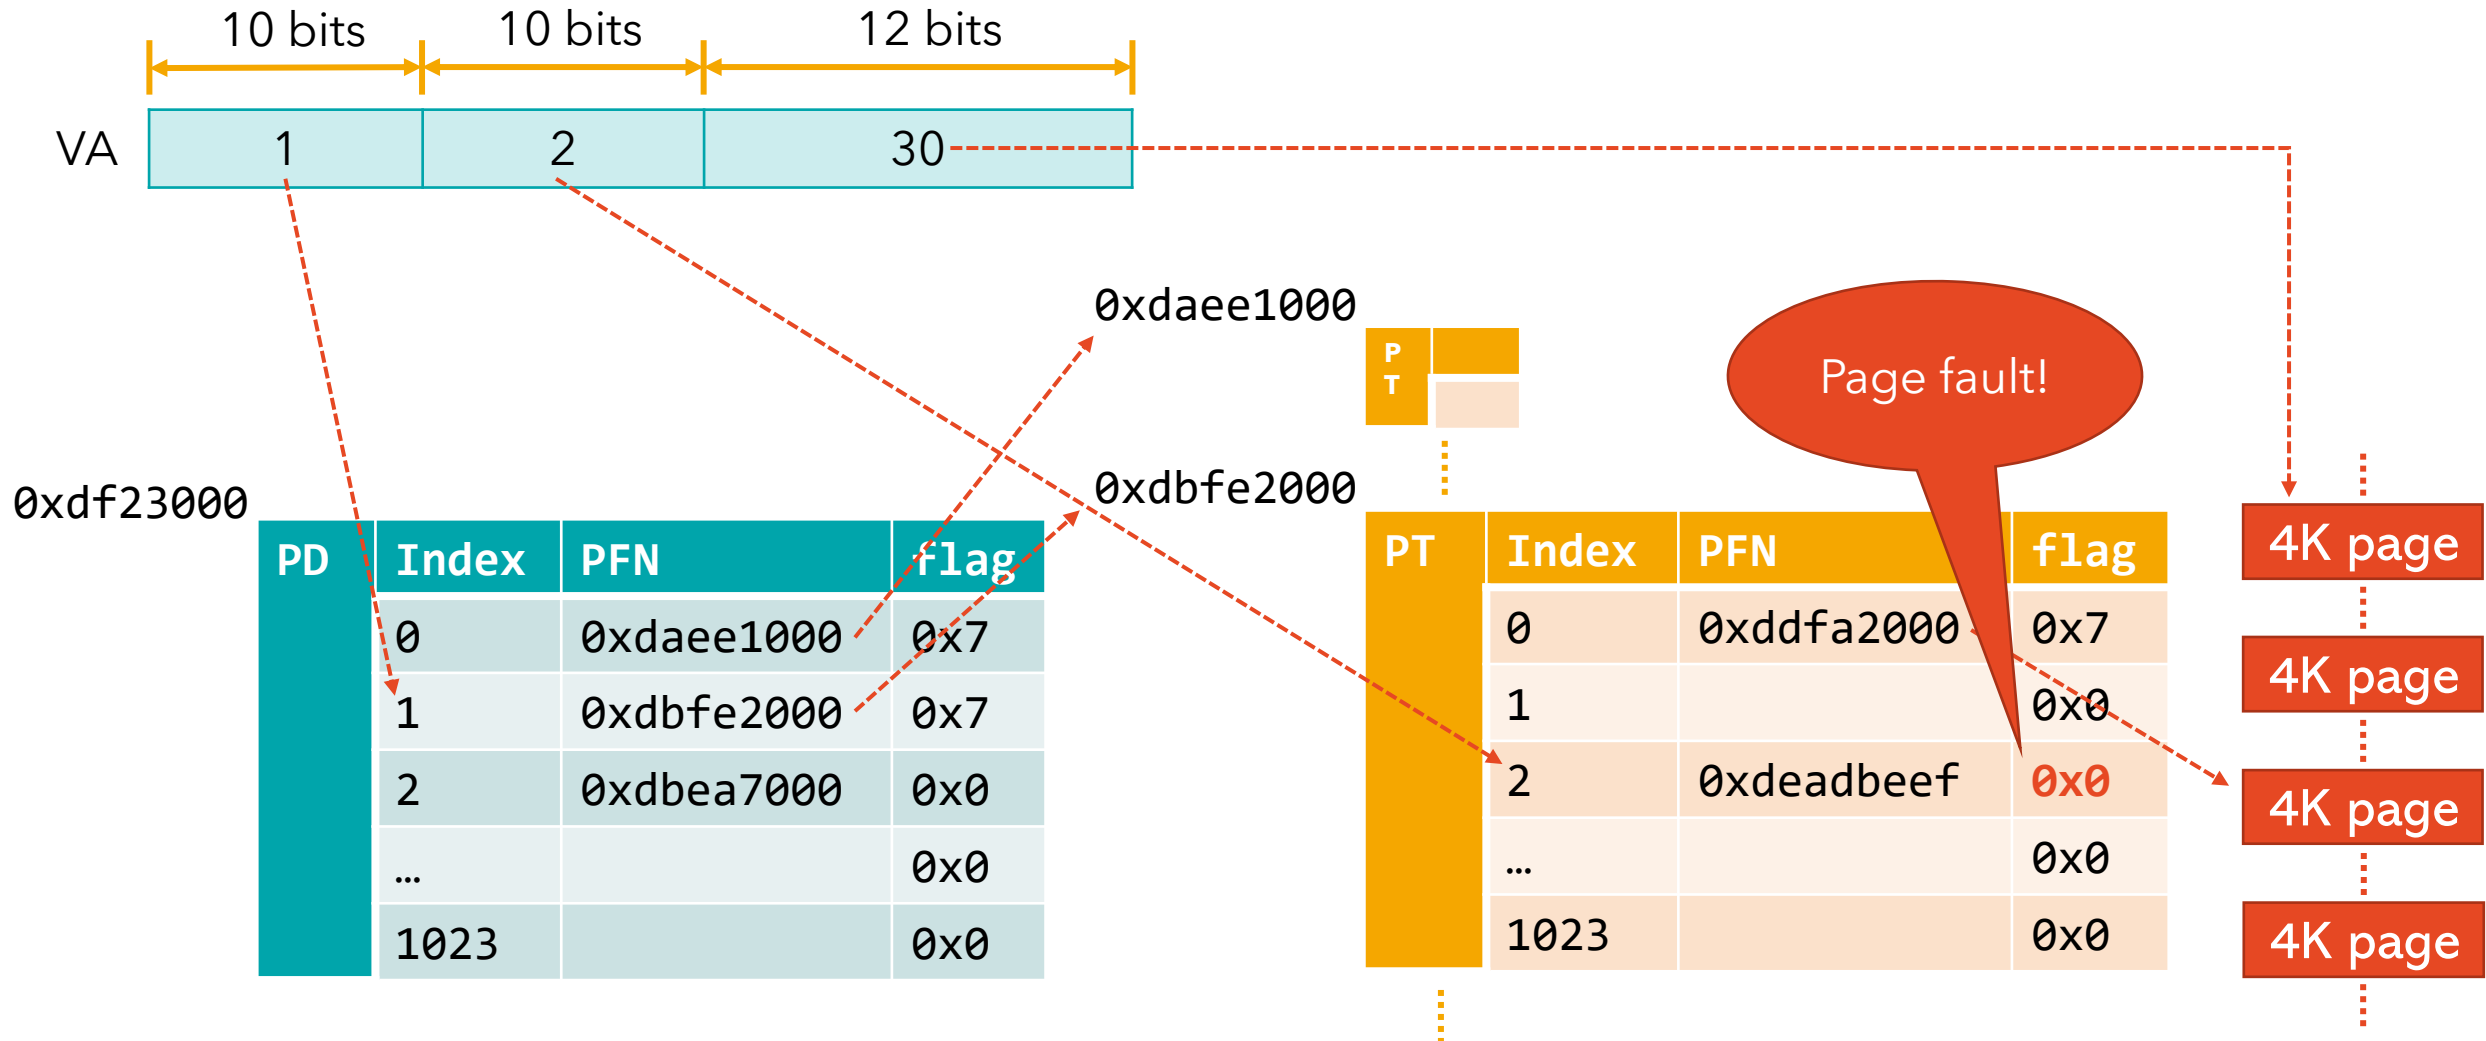


Error!

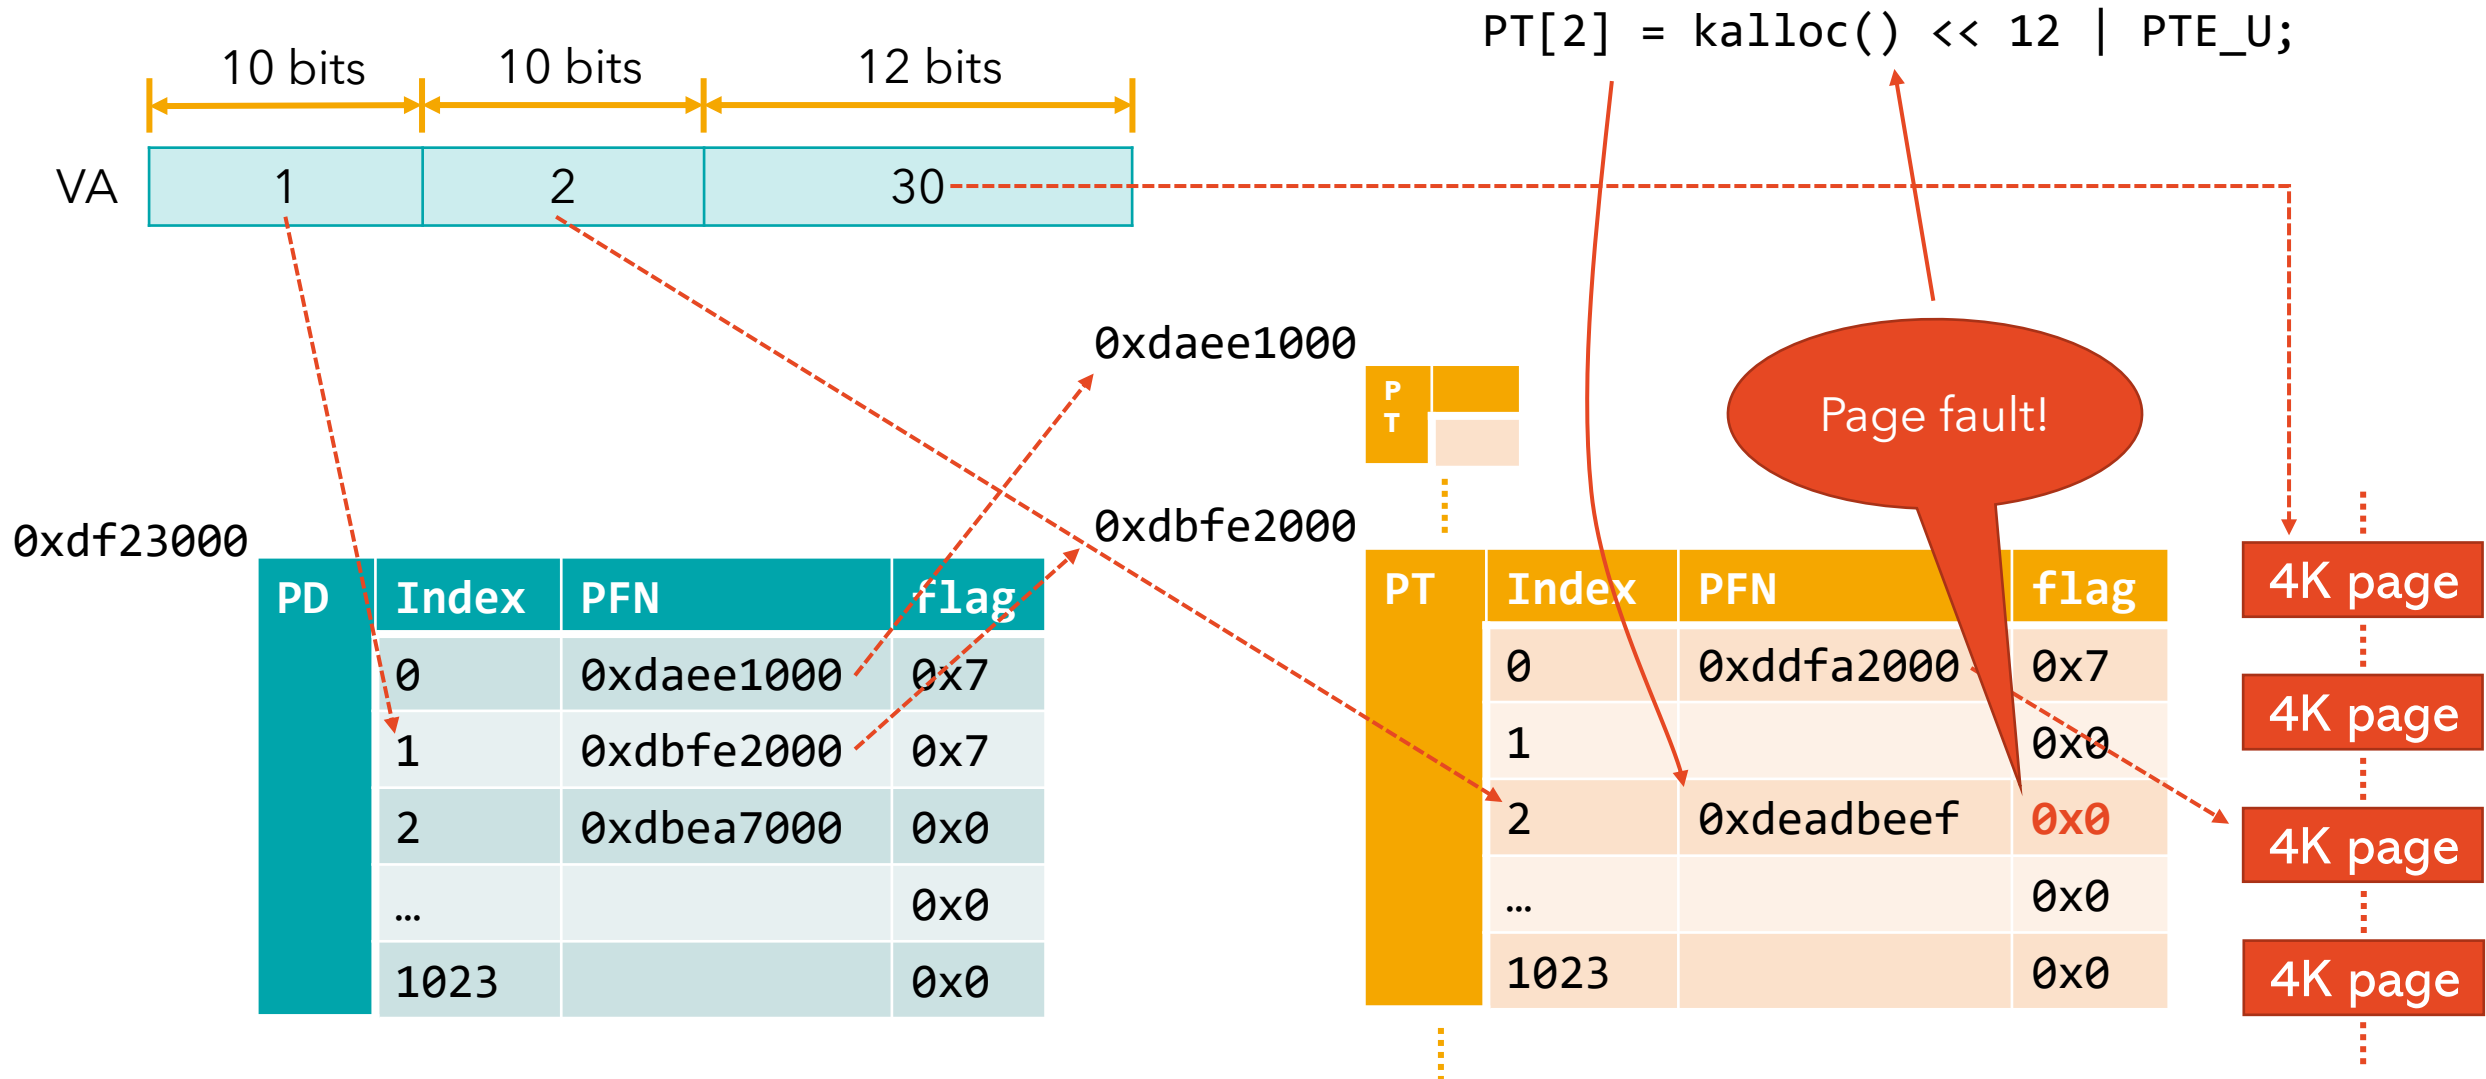
How can we implement custom page addressing within software?



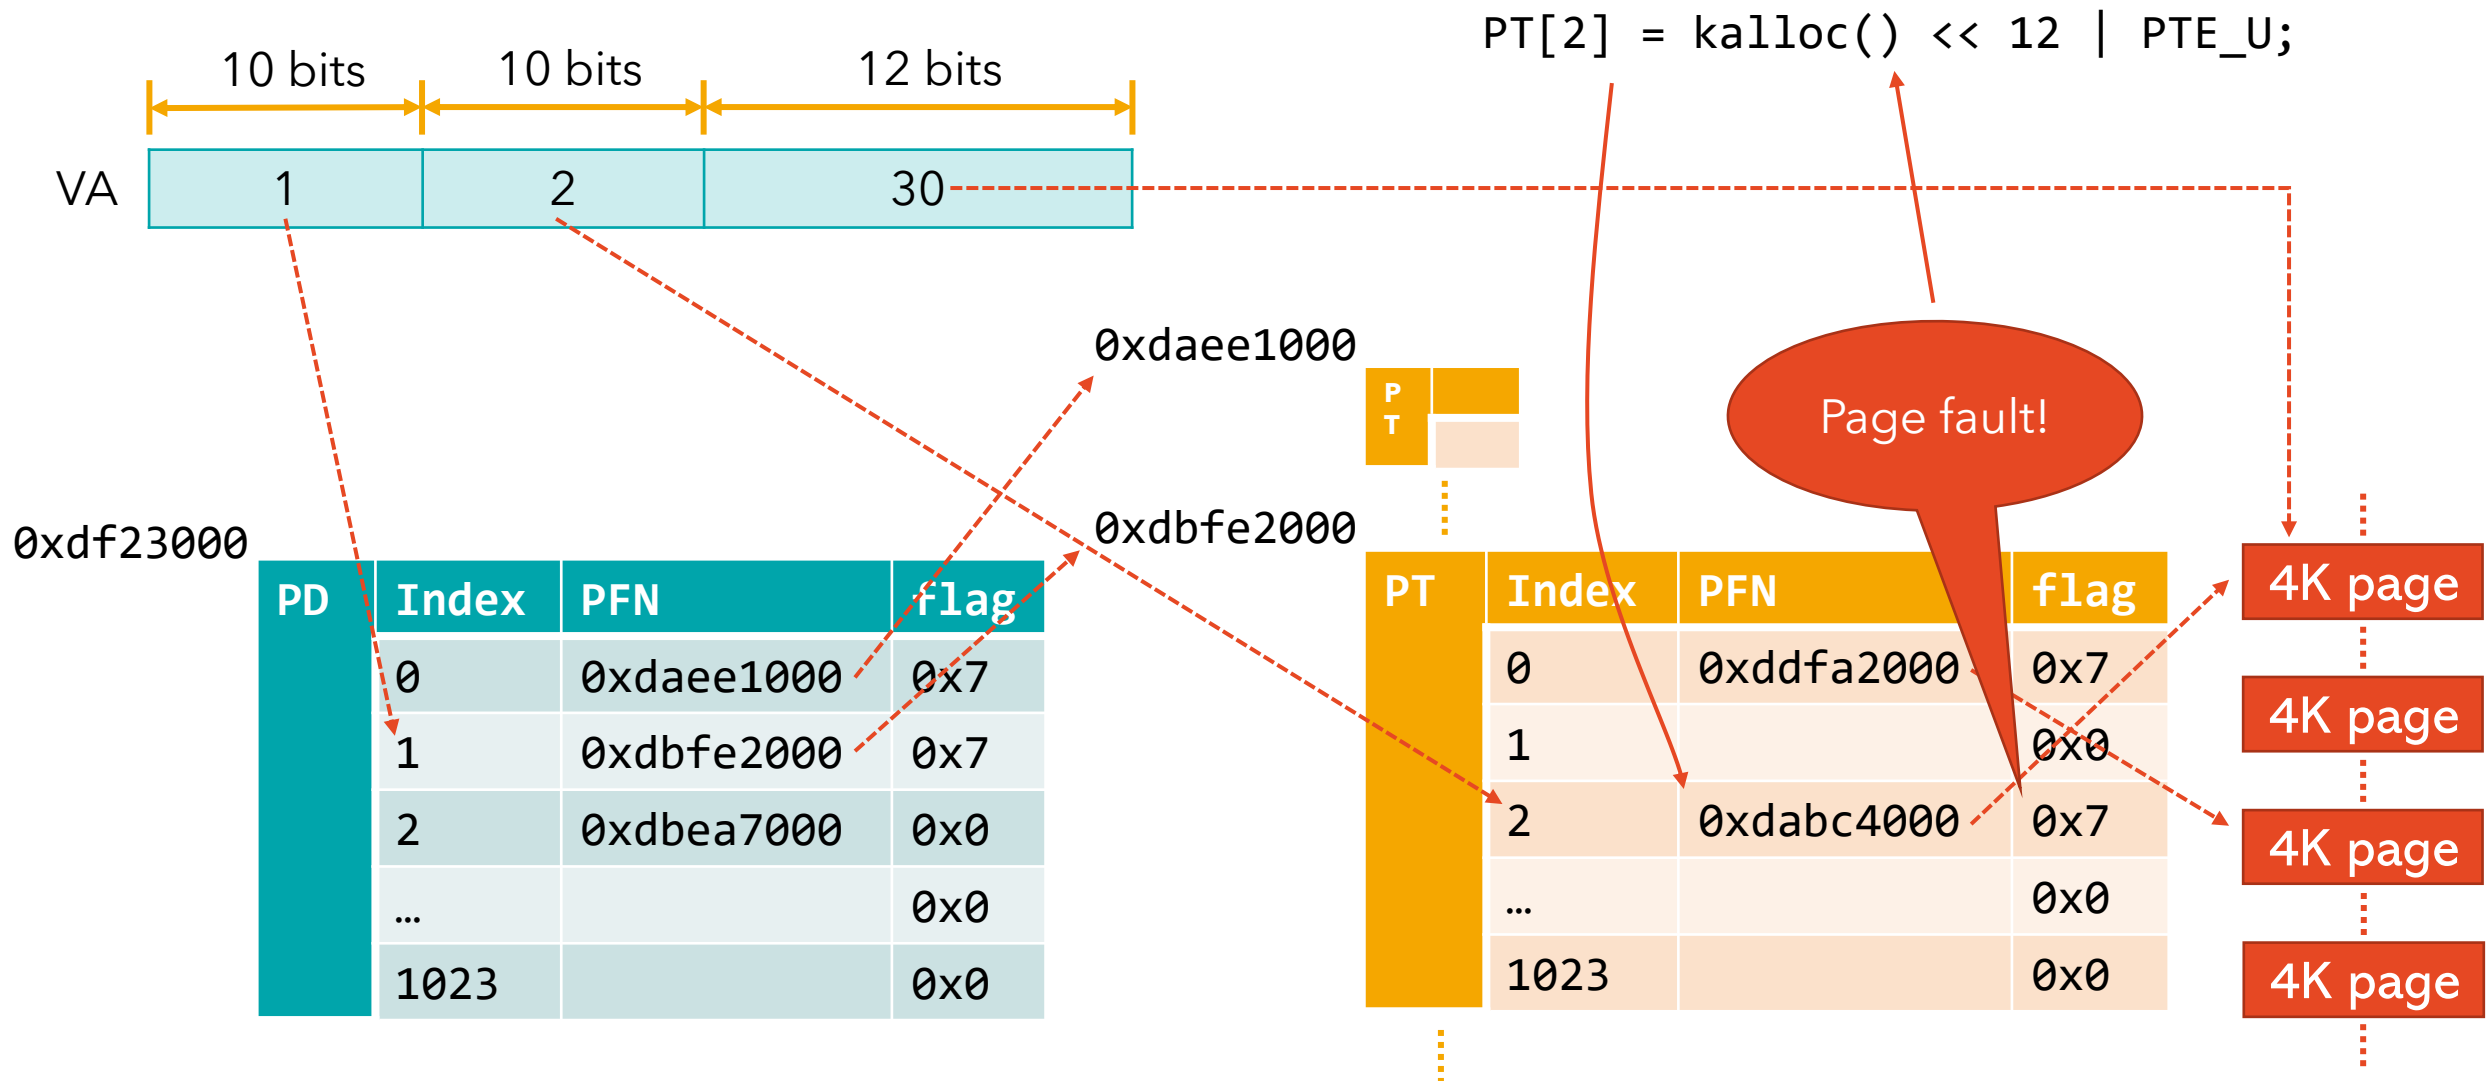
Page fault (I - occurs when PTE_P flag is not set)




Page fault (2 - kernel's page fault handler is executed)




Page fault (3 - code execution is resumed from page fault handler)



Triggering page fault on every page access allows the software to fully redirect VA!



How?



00001 00100
(1, 4)

Page fault!

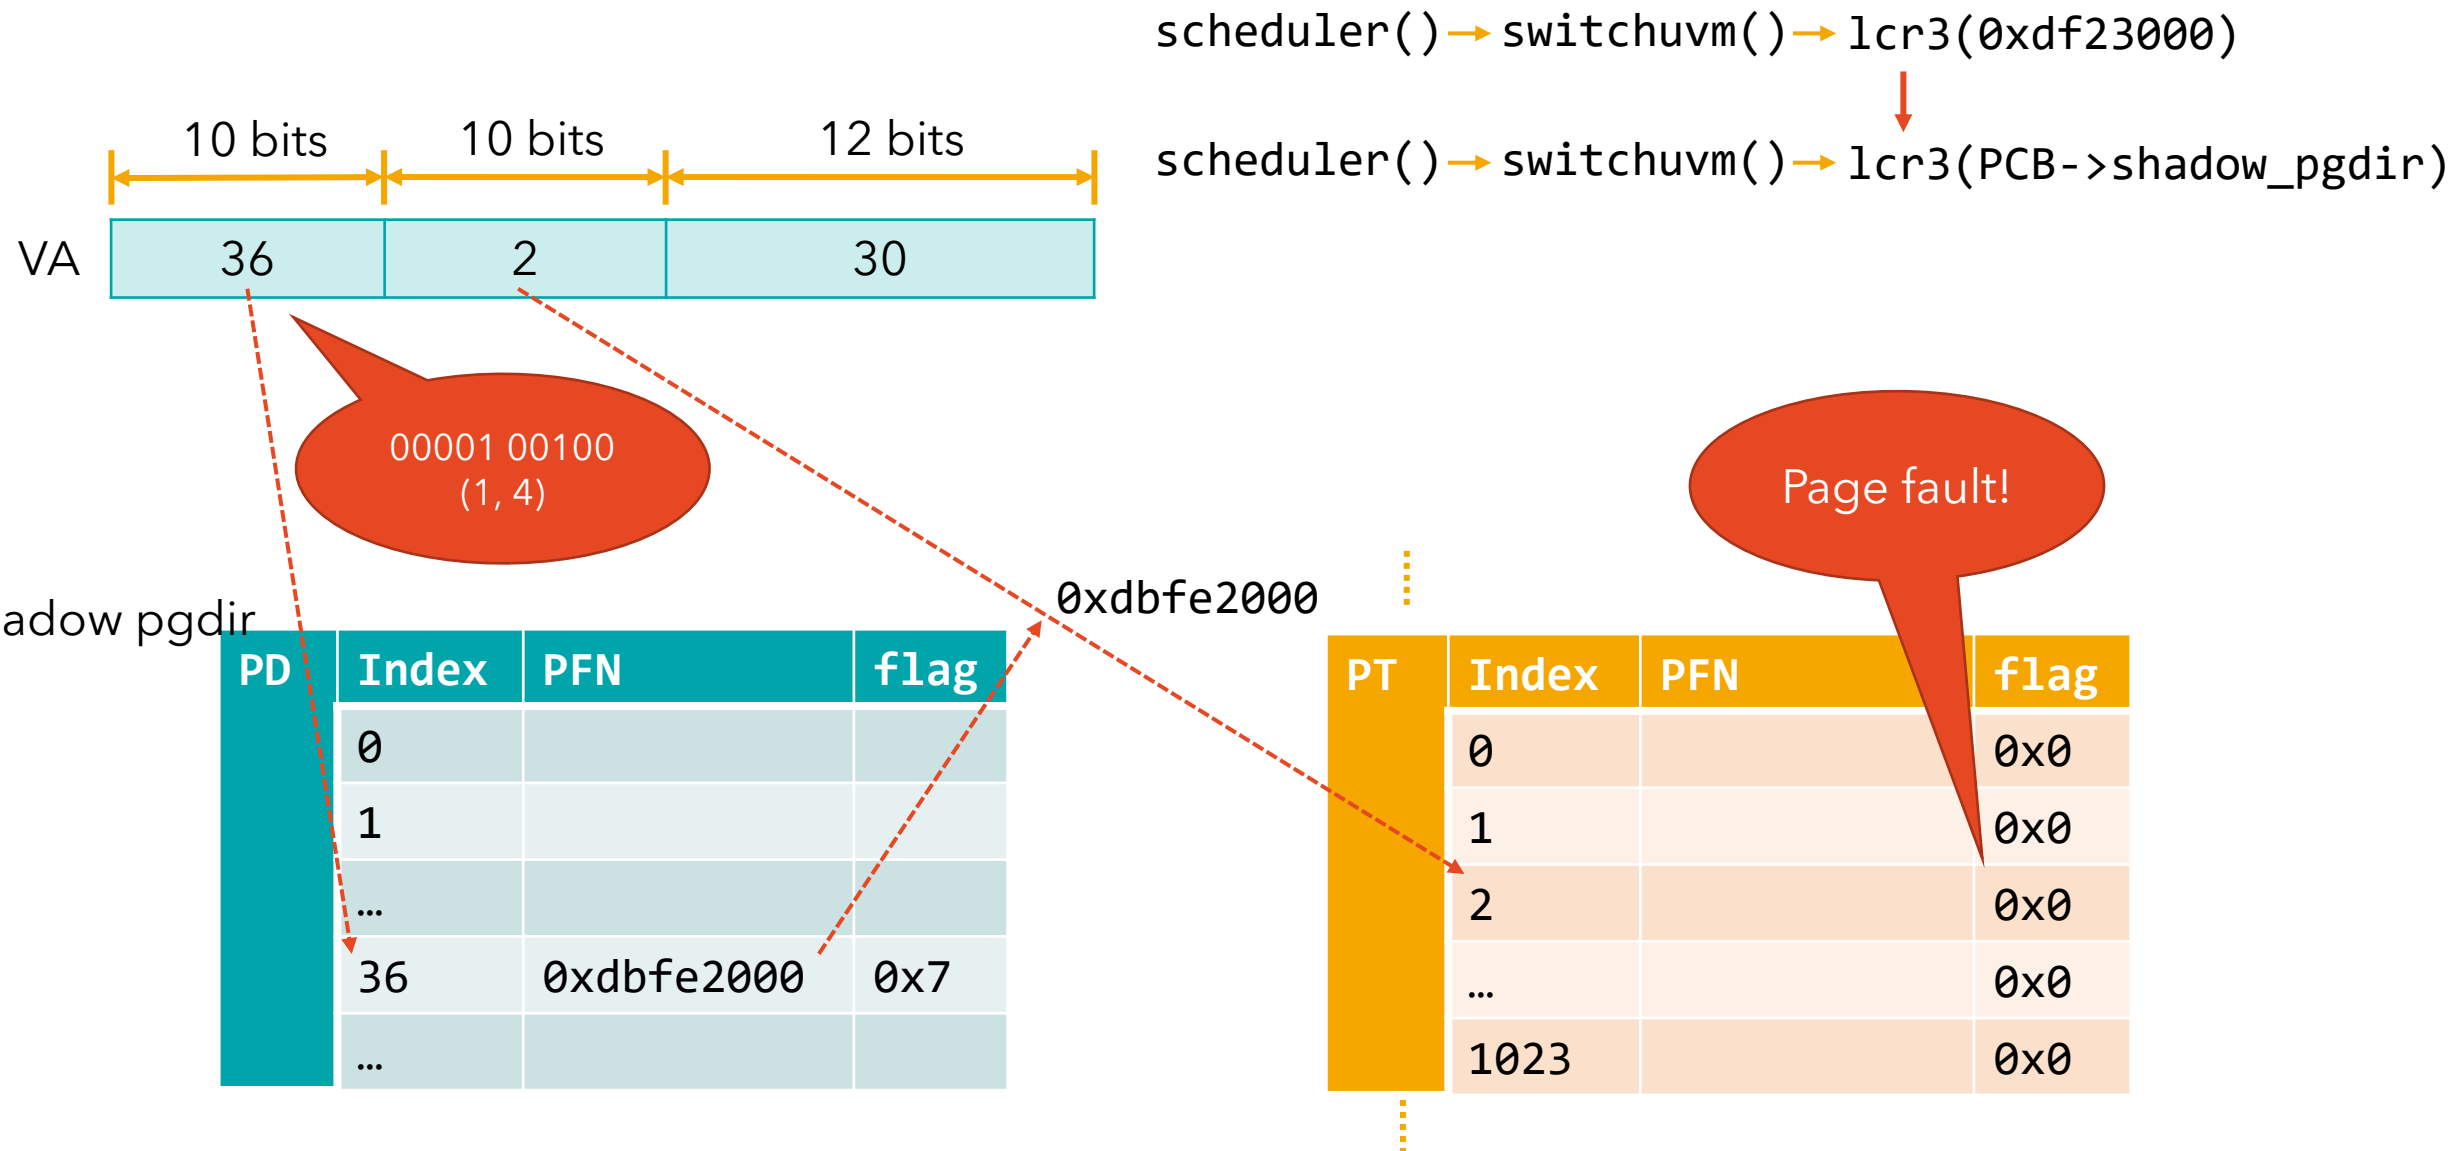
Shadow pgdir

PD	Index	PFN	flag
	0		
	1		
	...		
	36		0x0
	...		

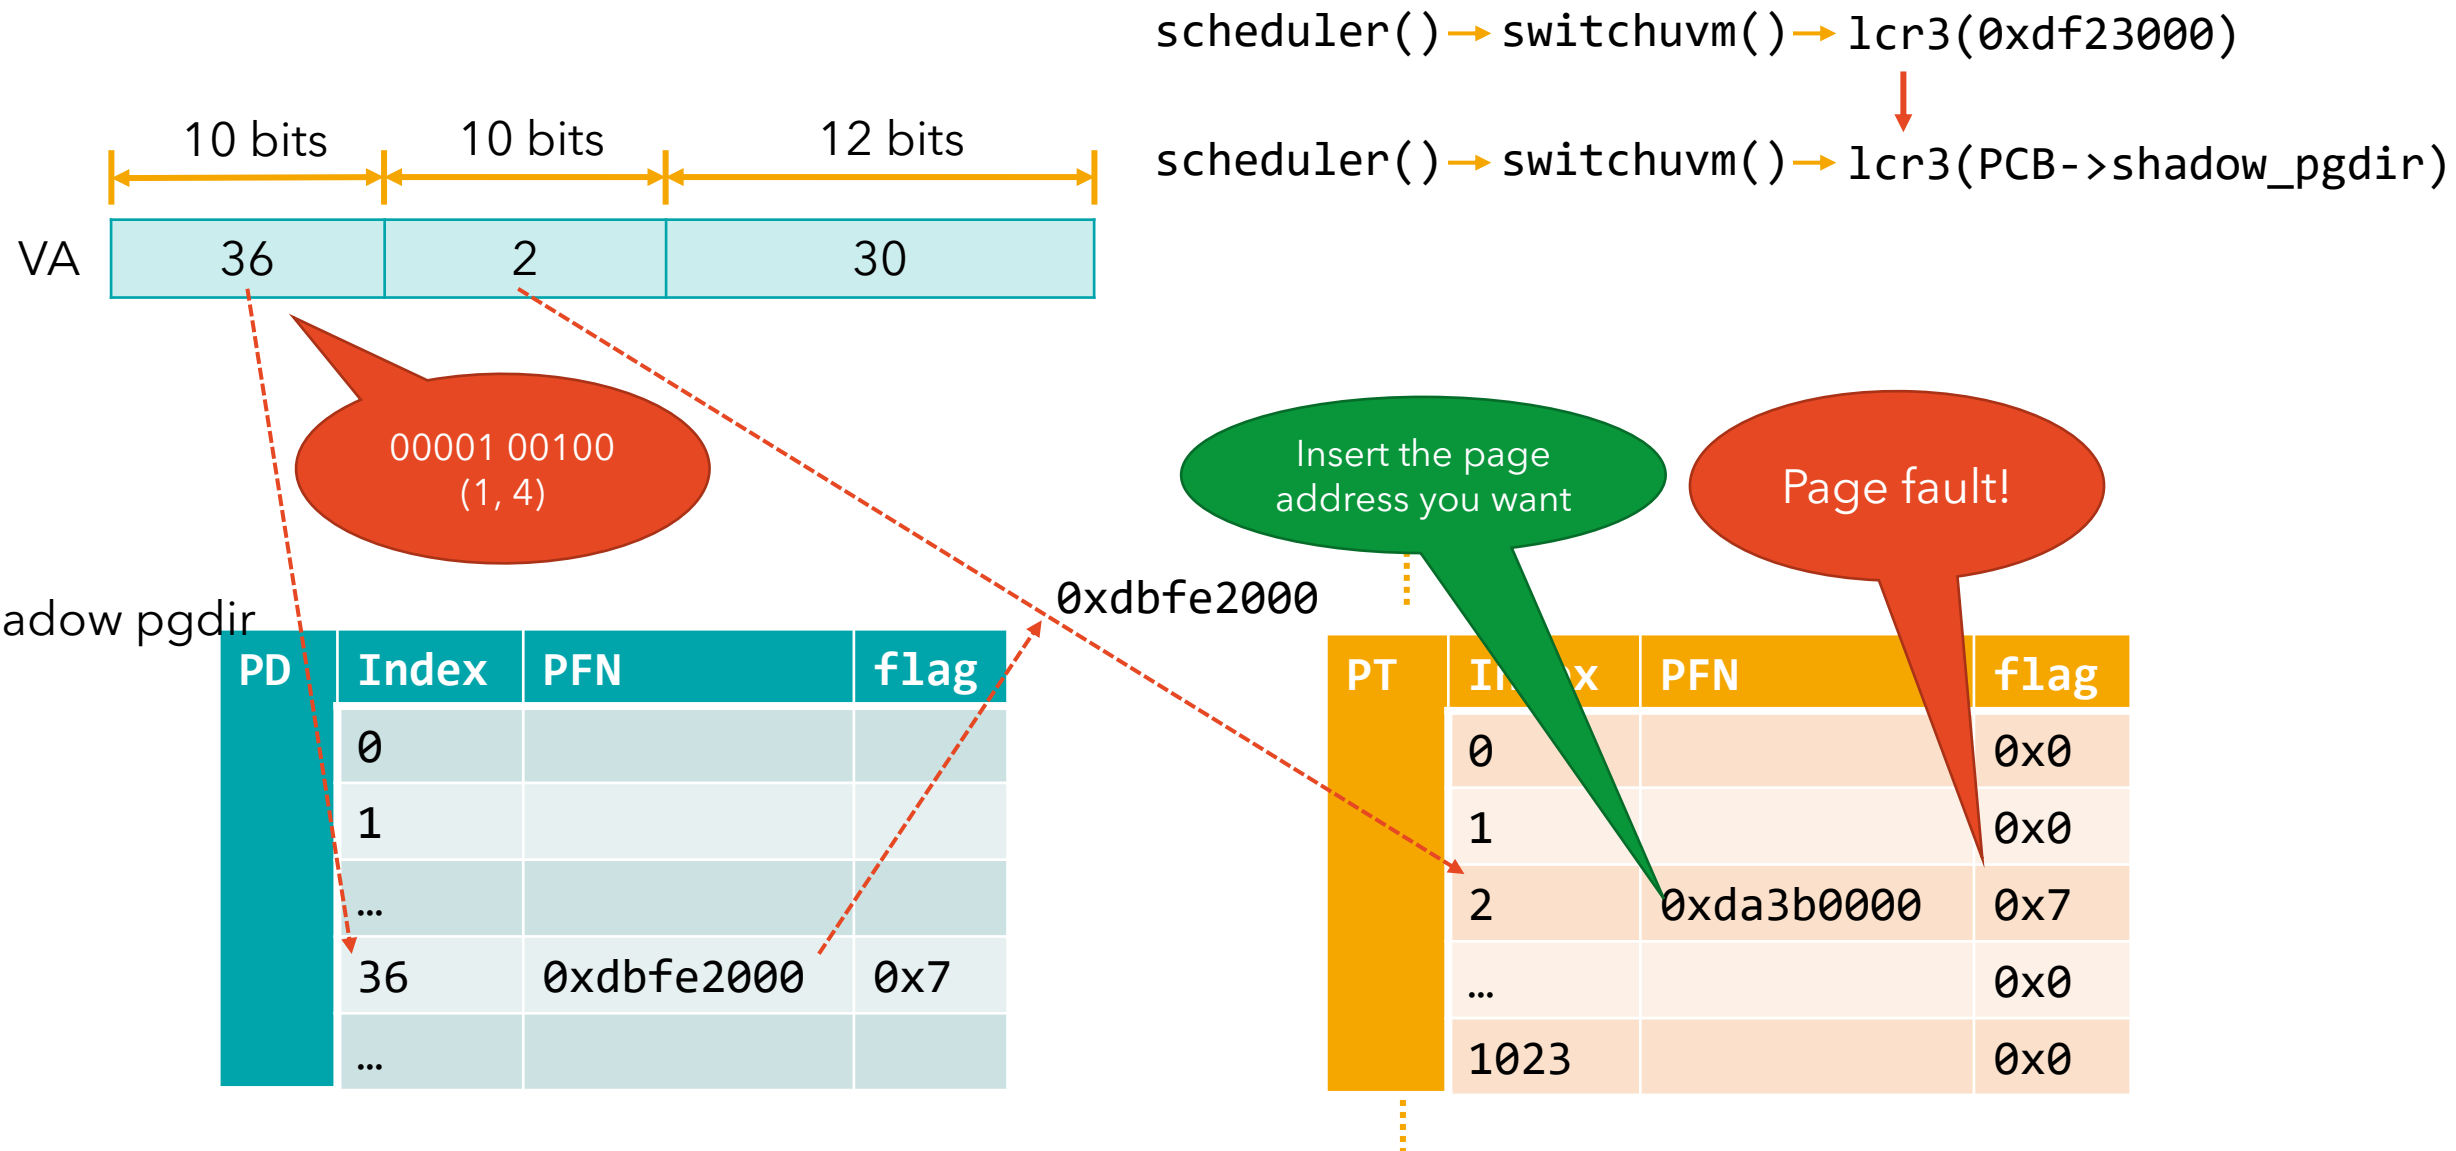
`scheduler() → switchvm() → lcr3(0xdf23000)`

`scheduler() → switchvm() → lcr3(PCB->shadow_pgdir)`

How?



How?





Software TLB

- Trigger a page fault on every new page access
- Reference the vanilla pgdir and make a copy to shadow_pgdir
- We maintain 2-level paging in this project but the CPU's MMU must never access PCB->pgdir
 - PCB->shadow_pgdir must return the appropriate PA instead

Trigger a page fault on every new page access (1/2)

- Setting the CR3 register tells the hardware to use it as the current page directory address
 - When is switchvm() called?
 - Can we abuse this to force page faults?
 - Hint: a new empty pgdir can be allocated with setupkvm();

```
vm.c:

// Switch TSS and h/w page table to correspond to process p.
void
switchvm(struct proc *p)
{
    ...

    ltr(SEG_TSS << 3);
    lcr3(V2P(p->pgdir)); // switch to process's address space
    popcli();
}
```


Reference the vanilla pgdir and make a copy to shadow_pgdir (2/2)

- The CPU's MMU must never access PCB->pgdir
 - It should use PCB->shadow_pgdir
- You must copy PA from PCB->pgdir and set it up to PCB->shadow_pgdir
 - Recall what walkpgdir() function did in vm.c

```
void pagefault(void)
{
    ...

    // Map pgdir's page address to shadow_pgdir's page table
    // XXX

    ...
}
```



Counting page accesses

- Once shadow_pgdir is set, page fault handler no longer runs if it was already handled before
- To count page accesses, shadow_pgdir must be cleared
- shadow_pgdir may contain addresses to page table itself
- Clearing all of shadow_pgdir will cause recursion loop during page fault

Handling sbrk'ed memory separately

- We use a reserved bit available on PDE

Page Directory Entry (4 MB)

31	...	22	21	20	...	13	12	11	...	9	8	7	6	5	4	3	2	1	0
Bits 31-22 of address				R S V D (0)	Bits 39-32 of address				P A T	AVL	G	P S (1)	D	A	P C D	P W T	U / S	R / W	P

Page Directory Entry

31	...	12	11	...	8	7	6	5	4	3	2	1	0	
Bits 31-12 of address						AVL	P S (0)	A V L	A	P C D	P W T	U / S	R / W	P

P: Present	D: Dirty
R/W: Read/Write	PS: Page Size
U/S: User/Supervisor	G: Global
PWT: Write-Through	AVL: Available
PCD: Cache Disable	PAT: Page Attribute Table
A: Accessed	

```
#define PTE_SBRK (1 << 10) // sbrk'ed memory (malloc)
```

Handling sbrk'ed memory separately

- We use a reserved bit available on PDE
- malloc => sbrk => growproc => PTE_SBRK set!
- pagefault => proc->last_va => if PTE_SBRK is set, clear!
- This allows page fault handler to run again on the same memory address
 - Which in turn allows us to “count” its occurrence
- <https://github.com/dgist-datalab/xv6/commit/27ce2ad>

Tips

- Try to understand 2-level paging fully before writing actual code
- Be careful of pointer accesses
- Reference mmu.h for helper macros
 - P2V, PDX, PTX, PTE_ADDR, ...
- Have a look at skeleton code's commit history
 - <https://github.com/dgist-datalab/xv6/commits/page-counter>
- `cprintf()` is buggy when called from page fault handler
 - You can set `#define LOG 1` from `vm.c` to forcefully enable `clprintf()`
 - If a page fault is triggered during `cprintf()`...
 - It'll cause a crash during normal operations

```
pagefault++
Page fault by process "init" (pid: 1) at 0x38a
virt_to_phys: translated "init"(1)'s VA 0x38a to PA 0xdf3838a (pgdir)
Clearing last_pde_entry (0x8dffe000)
virt_to_phys: translated "init"(1)'s VA 0x38a to PA 0xdf3838a (shadow_pgdir)
pagefault--
pagefault++
Page fault by process "init" (pid: 1) at 0x38a
virt_to_phys: translated "init"(1)'s VA 0x38a to PA 0xdf3838a (pgdir)
Clearing last_pde_entry (0x8dffe000)
virt_to_phys: translated "init"(1)'s VA 0x38a to PA 0xdf3838a (shadow_pgdir)
pagefault--
lapicid 0: panic: acquire
801046aa 801007ed 80107203 80105b65 8010588f 80101170 80104e29 80104ad9 80105ab5 8010588f
```

In action

- #define LOG 1 (vm.c)

```
x - □ arter97@arter97-x1: ~/lab/os/xv6
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
pagefault++
Page fault by process "initcode" (pid: 1) at 0x0
virt_to_phys: translated "initcode"(1)'s VA 0x0 to PA 0xdf7c000 (pgdir)
→ Allocated pgtable at 0x8df7a000
→ virt_to_phys: translated "initcode"(1)'s VA 0x0 to PA 0xdf7c000 (shadow_pgdir)
pagefault--
pagefault++
Page fault by process "initcode" (pid: 1) at 0x21
virt_to_phys: translated "initcode"(1)'s VA 0x21 to PA 0xdf7c021 (pgdir)
→ Clearing last_pde_entry (0x8df7a000)
→ virt_to_phys: translated "initcode"(1)'s VA 0x21 to PA 0xdf7c021 (shadow_pgdir)
pagefault--
pagefault++
Page fault by process "initcode" (pid: 1) at 0x1c
virt_to_phys: translated "initcode"(1)'s VA 0x1c to PA 0xdf7c01c (pgdir)
→ Clearing last_pde_entry (0x8df7a000)
→ virt_to_phys: translated "initcode"(1)'s VA 0x1c to PA 0xdf7c01c (shadow_pgdir)
pagefault--
```




Extra points

- If you found a mistake in the skeleton code
- If you found a different/better method
- Kernel memory leak test has been disabled in usertests
 - <https://github.com/dgist-datalab/xv6/commit/63776c9cc>
 - If you manage to pass this test (after reverting the commit)
- Write a report for extra points



Finally ...

Do NOT hesitate to ask questions!



Thank You!